

An Introduction to R for the Geosciences: Regression

Gavin Simpson

Institute of Environmental Change & Society
and
Department of Biology
University of Regina

30th April — 3rd May 2013

Simple linear regression

- **Simple linear regression** is a statistical model that assumes a linear relationship between a continuous response variable y and one or more, usually continuous, predictor variables, $X = x_1, \dots, x_n$
- Three major purposes of such models
 - ▶ to describe the linear relationship between y and X
 - ▶ to determine how much variation (uncertainty) in y can be explained by the relationship with X , and
 - ▶ to predict new values of y from new values of X
- A linear model is linear in its parameters only — the fitted response can be non-linear in the English sense of the word

Simple linear regression

- Consider first the case of a single predictor variable x and its relationship to y
- A suitable form for such a model is



$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

- We need to estimate two parameters (β_0 and β_1)
- β_0 is the intercept, the mean of the probability distribution of y when x is 0
- β_1 is often called the slope, it measures the rate of change in y for a per unit change in x
- Estimate the parameters using least-squares, solving this model by minimising Residual Sum of Squares

$$RSS = \sum_{i=1}^n \varepsilon_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

Least-squares


- $\hat{\beta}_0$ is the estimate of the intercept
- $\hat{\beta}_1$ is the estimate of the slope
- Observed points (y_i) are open circles
- Fitted points (\hat{y}_i) are filled circles
- Fitted model/line is solid line through \hat{y}_i
- Dashed lines between y_i and \hat{y}_i are the residuals (ε_i)
- Thick black line shows prediction of \hat{y}_{new} given a new x value of 2.5



Least-squares

- Data generated from true mean function
- Least squares estimates of mean function

$$\beta_0 = 0.7 \quad \beta_1 = 0.8 \quad \hat{\beta}_0 = 0.3201 \quad \hat{\beta}_1 = 0.9987$$



Least-squares

- Data are 20 observations generated from the following model

$$y_i = 0.7 + 0.8x_i + \varepsilon_i$$

$$\varepsilon_i \sim N(\mu = 0, \sigma = 1)$$

- Fitted model gives $\hat{\beta}_0 = 0.3201$ and $\hat{\beta}_1 = 0.9987$
- F-ratio for this fitted model is 66.43, which has a p -value of > 0.0001 from a F distribution with 1 and $n - 2$ (20) degrees of freedom
- This is equivalent of testing our model against the Null model (null hypothesis) that

$$y_i = \beta_0 + \varepsilon_i$$

- where β_0 is just the sample mean, \bar{y}_i , i.e. that there is no variation in y given x

Assumptions of least squares regression

- ① The linear model correctly describes the functional relationship between y and X
 - ▶ If violated the estimate of predictor variances (σ^2) will be inflated
 - ▶ Incorrect model specification can show itself as patterns in the residuals
- ② x_i are measured without error
 - ▶ Allows us to isolate the error component as random variation in y
 - ▶ Estimates $\hat{\beta}$ will be biased if there is error in X — often ignored!
- ③ For any given value of x_i , the sampled y_i values are independent with normally distributed errors
 - ▶ Independence and normality of errors allows us to use parametric theory for confidence intervals and hypothesis tests on the F-ratio.
- ④ Variances are constant along the regression line/model
 - ▶ Allows a single constant variance σ^2 for the variance of the regression line/model
 - ▶ Non-constant variances can be recognised through plots of residuals (amongst others) — i.e. residuals get wider as the values of y increase.

Fitting linear models in R

Typical model call & output from R. Next few slides explain the salient results

```
> mod <- lm(Age ~ Depth, data = agedat)
> summary(mod)
```

Call:

```
lm(formula = Age ~ Depth, data = agedat)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.3808	-7.7115	0.7053	6.1577	16.7818

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	21.2480	3.5626	5.964	2.02e-06 ***
Depth	5.5760	0.3208	17.384	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.131 on 28 degrees of freedom

Multiple R-squared: 0.9152, Adjusted R-squared: 0.9122

F-statistic: 302.2 on 1 and 28 DF, p-value: < 2.2e-16

Fitting linear models in R

- Estimate is β_j , the model coefficients, on log scale (base e)
- For 1m increase in sediment Depth, sediment Age decreases by 5.576kyrs
- t -value is the t statistic, the ratio of the estimate and its standard error $t = \frac{\hat{\beta}_j}{\hat{se}_j}$
- p -value is probability of achieving a t as large or larger than the one observed under null hypothesis
- Intercept of interest — sediment age at 0m sediment depth

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	21.2480	3.5626	5.964	2.02e-06 ***
Depth	5.5760	0.3208	17.384	< 2e-16 ***

Fitting linear models in R

- Residual standard deviation $\hat{\sigma} = 9.131$; a measure of the variance of the residuals
- r^2 is the coefficient of determination, the ratio of the variance explained to the total variance; a measure of how much variance is explained

$$r^2 = \frac{SS_{\text{regression}}}{SS_{\text{regression}} + RSS} = 1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}}$$

- Adjusted r^2 takes into account number of predictors in the model

$$r^2_{\text{adj}} = 1 - \frac{SS_{\text{residual}}/[n - (p + 1)]}{SS_{\text{total}}/(n - 1)}$$

- If we added a redundant predictor to model r^2 would increase. r^2_{adj} attempts to control for this phenomenon

Residual standard error: 9.131 on 28 degrees of freedom

Multiple R-squared: 0.9152, Adjusted R-squared: 0.9122

F-statistic: 302.2 on 1 and 28 DF, p-value: < 2.2e-16

Fitting linear models in R

- F is the F -ratio, the ratio of the regression and residual variances (Mean squares)

$$F = \frac{\sum_{i=1}^n (\hat{y}_i - \bar{y})^2 / p}{\sum_{i=1}^n (y_i - \hat{y}_i)^2 / [n - (p + 1)]} = \frac{\text{MS}_{\text{regression}}}{\text{MS}_{\text{residual}}}$$

- Probability of F greater than or equal to observed from F -distribution with p and $n - (p + 1)$ degrees of freedom

```
> anova(mod)
```

Analysis of Variance Table

Response: Age

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Depth	1	25195.9	25195.9	302.2	< 2.2e-16 ***
Residuals	28	2334.5	83.4		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Fitting linear models in R — Hypothesis testing

- t tests are tests the H_0 that $\hat{\beta}_j = 0$
- F tests the ratio of variance explained to unexplained
- With single predictor, t test for length and F of model are equivalent
- More generally we can think of F as comparing

$$y_i = \beta_0 + \varepsilon_i$$

with

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

```
> mod0 <- lm(Age ~ 1, data = agedat)
> anova(mod0, mod) ## same as anova(mod)
```

Analysis of Variance Table

Model 1: Age ~ 1

Model 2: Age ~ Depth

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	29	27530.4				
2	28	2334.5	1	25196	302.2	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R's model formula

- R uses a slightly modified version of the Wilkinson-Rogers (Wilkinson & Rogers (1973; Applied Statistics **22**;392–399) notation to symbolically describe statistical models

```
mod <- lm(Y ~ x1 + x2, data = mydata)
```

- Intercept implied; suppress with `- 1` or `+ 0`

```
mod <- lm(Y ~ x1 + x2 - 1, data = mydata)
```

- Interaction terms with `a : b`

```
mod <- lm(Y ~ x1 + x2 + x1:x2, data = mydata)
```

- Can be simplified using `a * b`

```
mod <- lm(Y ~ x1 * x2, data = mydata)
```

- Shortcut to add all variables to model is `.` (Careful!)

```
mod <- lm(Y ~ ., data = mydata)
```

- Polynomials via `I(x^2)` or `poly(x, 2)`

```
mod <- lm(Y ~ x + I(x^2), data = mydata)
```

```
mod <- lm(Y ~ poly(x, 2), data = mydata)
```

R's model formula

- Note the use of the `data` argument. This is a data frame (or list) containing the variables to include in the model

```
mod <- lm(Y ~ x1 + x2, data = mydata)
```
- You **never** want to do this

```
mod <- lm(mydata$Y ~ mydata$x1 + mydata$x2)
```
- Apart from taking longer to type, the `predict()` method won't work easily
- Can include functions in formula as `poly(x, 2)` earlier

```
mod <- lm(Y ~ log(x), data = mydata)
```
- Better to do this if you can than transform data & store both transformed and untransformed variable in your data frame. Not least because `predict()` just works
- You can exclude variables too

```
mod <- lm(Y ~ x1 + x2 - x3, data = mydata)
```

Multiple regression

- The simple regression model readily generalises to the situation where we have m predictors not just one.

$$y_i = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_m x_m + \varepsilon$$

- Now we have $m + 1$ parameters to estimate, one for intercept and one each for the m predictors x_m
- It is tedious to write all that out, so we collect the β_m into a vector β and all the predictors (including the intercept, a vector of 1s) into the model matrix, X , then rewrite the model as

$$y = \mathbf{X}\beta + \varepsilon$$

Akaike information criterion

- Akaike information criterion (AIC) is an index of fit that takes account of the parsimony of the model by penalising for the number of parameters
- The more parameters in the model the better the fit — if you have as many parameters as data points then the fit is perfect but the model has no explanatory power! A Trade-off.
- AIC is useful as it explicitly penalises any superfluous parameters in the model by adding $2p$ where p is the number of parameters to the variance or deviance of the model.

$$\text{AIC} = -2 \times \max \log\text{likelihood} + 2p$$

- Associated is Bayes information criterion (BIC), which applies a stronger penalty of $p \log n$, where n is number of observations
- For linear regression the $-2 \times \max \log\text{likelihood}$ is $n \log(RSS/n) + \text{constant}$, where RSS is the residual sums of squares.

Akaike information criterion

- We use AIC and BIC to compare two or more **nested** models
- **Nested** means that one model is a subset of the other
- The model with the smallest AIC or BIC is to be preferred
- Note that you can get negative values for AIC and BIC. This is fine, just go for the smallest value: e.g. -21.5 is better than -15.4
- Difference in AIC of 2 is expected with a redundant parameter
- Models with AIC differing by 2 or less are effectively the same
- **AIC()** & **BIC()** methods can be used to extract IC from fitted model objects

```
> AIC(mod)
[1] 221.7669
> BIC(mod)
[1] 225.9705
```

ANOVA — the Analysis of Variance

- **ANOVA** is a general statistical technique for partitioning and analysing the variation in a continuous response variable
- Earlier we used ANOVA to partition the variance in a response variable into components explained by explanatory variables and a residual component not explained by the regression model
- A slightly more restricted view of ANOVA is that it is a technique for partitioning the variation in a response variable into that explained or unexplained by one or more categorical predictor variables or factors
- The categories of each factor are the groups or experimental treatments
- Often the focus is on comparing the mean of the response variable between groups
- We won't dwell too much on the distinction between regression and ANOVA — they are effectively the same and in R we use the same fitting function, e.g. `lm()`

Simple one-way ANOVA

- One-way ANOVA designs deal with only a single factor or predictor variable
- The single factor comprises 2 or more groups
- Medley & Clements (1998) studied the response of diatom communities to heavy metals (esp. Zinc, Zn) in streams in the Rocky Mountain region of Colorado, USA
- They sampled a number of stations (4–7) on six streams known to be polluted by heavy metals
- Several variables were measured at each station, inc. Zn concentration, diatom species richness and diversity, and proportion of diatom cells belonging to the diatom *Achnanthes minutissima*
- Zn concentration used to group sites into four categories;
- Is there a difference in species diversity between the four Zn categories?

Rocky mountain diatoms

```
> diatom <- read.csv("medley.csv")
> diatom$ZINC <- factor(diatom$ZINC, levels = c("BACK","LOW","MED","HIGH"))
> ## Drop some superfluous columns
> diatom <- diatom[, 1:3]
> head(diatom)
  STREAM ZINC DIVERSITY
1  Eagle BACK      2.27
2  Eagle HIGH      1.25
3  Eagle HIGH      1.15
4  Eagle  MED      1.62
5   Blue BACK      1.70
6   Blue HIGH      0.63
> str(diatom)
'data.frame':   34 obs. of  3 variables:
 $ STREAM   : Factor w/ 6 levels "Arkan","Blue",...: 4 4 4 4 2 2 2 2 2 ...
 $ ZINC      : Factor w/ 4 levels "BACK","LOW","MED",...: 1 4 4 3 1 4 1 1 4 3 ...
 $ DIVERSITY: num 2.27 1.25 1.15 1.62 1.7 0.63 2.05 1.98 1.04 2.19 ...
> table(diatom$ZINC)
```

```
BACK  LOW  MED HIGH
   8    8    9    9
```

```
> table(diatom$STREAM)
```


```
Arkan  Blue Chalk Eagle Snake Splat
   7    7    5    4    5    6
```

```
> with(diatom, table(ZINC, STREAM))
```

```
  STREAM
ZINC   Arkan Blue Chalk Eagle Snake Splat
BACK    0    3    0    1    1    3
LOW     5    0    2    0    0    1
MED     2    2    1    1    1    2
HIGH    0    2    2    2    3    0
```

Rocky Mountain diatoms

```
> boxplot(DIVERSITY ~ ZINC, data = diatom)
```



Rocky Mountain diatoms — ANOVA

```
> zn.lm1 <- lm(DIVERSITY ~ ZINC, data = diatom)
> summary(zn.lm1)
```

```
Call:
lm(formula = DIVERSITY ~ ZINC, data = diatom)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.03750	-0.22896	0.07986	0.33222	0.79750

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.79750	0.16478	10.909	5.81e-12 ***
ZINCLOW	0.23500	0.23303	1.008	0.3213
ZINCMED	-0.07972	0.22647	-0.352	0.7273
ZINCHIGH	-0.51972	0.22647	-2.295	0.0289 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4661 on 30 degrees of freedom
Multiple R-squared: 0.2826, Adjusted R-squared: 0.2108
F-statistic: 3.939 on 3 and 30 DF, p-value: 0.01756

```
> anova(zn.lm1)
```

Analysis of Variance Table

Response: DIVERSITY

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
ZINC	3	2.5666	0.8555	3.9387	0.01756 *
Residuals	30	6.5164	0.2172		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Rocky Mountain diatoms — different parametrisation

- Previous model (`zn.lm1`) contained an intercept
- To maintain identifiability, need to set one level of ZINC as reference level and express model as differences in mean diversity from this reference level
- If we re-parametrise and drop the intercept, the estimates are the group means

```
> zn.lm0 <- lm(DIVERSITY ~ ZINC - 1, data = diatom)
> coef(zn.lm0)
ZINCBACK ZINCLOW ZINCMED ZINCHIGH
1.797500 2.032500 1.717778 1.277778
> with(diatom, aggregate(DIVERSITY, list(ZINC = ZINC), mean))
      ZINC      x
1 BACK 1.797500
2 LOW  2.032500
3 MED  1.717778
4 HIGH 1.277778
```

Rocky Mountain diatoms — dummy variable coding in R

- Both models use **Treatment** contrasts
- Normally, one level is set as baseline and dropped, and contrasts code so as to reflect differences in that level from reference level
- Other contrasts are available, such as Helmert contrasts, see `?contrasts`

```
> model.matrix(zn.lm1)
  (Intercept) ZINCLOW ZINCMED ZINCHIGH
1           1         0         0         0
2           1         0         0         1
3           1         0         0         1
4           1         0         1         0
5           1         0         0         0
6           1         0         0         1
7           1         0         0         0
8           1         0         0         0
9           1         0         0         1
10          1         0         1         0
....
attr("assign")
[1] 0 1 1 1
attr("contrasts")
attr("contrasts")$ZINC
[1] "contr.treatment"
```


Outliers

- **Outlier** — observation which is inconsistent with the rest of the observations in a sample.
- An observation can be an outlier due to the response variable(s) or one or more of the predictor variables having values outside their expected limits.
- Identify outliers at EDA stage for investigation and evaluation, *not* rejection and deletion.
- An outlier may result from
 - ▶ incorrect measurement,
 - ▶ incorrect data entry,
 - ▶ transcription error,
 - ▶ recording error,
- Outliers are model dependent
- Two main concepts
 - ▶ **Leverage** — Potential for an outlier to be influential
 - ▶ **Influence** — Observation is influential if its deletion substantially changes the results

Leverage measures

Projection or Hat matrix

$$\mathbf{H} = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T$$

where \mathbf{X} is the $n \times p$ matrix of x values, the **parameters** in the model.

\mathbf{H} is an $n \times n$ matrix.

Hat matrix

$$\mathbf{H} = \begin{bmatrix} h_{11} & h_{12} & \dots & h_{1n} \\ h_{21} & h_{22} & \dots & h_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ h_{n1} & h_{n2} & \dots & h_{nn} \end{bmatrix}$$

- Hat matrix is so called because it puts a **hat on** \mathbf{Y} : $\hat{\mathbf{Y}} = \mathbf{H}\mathbf{Y}$.
- **Leverage** of an observation i is denoted h_{ii} — the i th element of the diagonal of \mathbf{H} .
- Leverage ranges from $1/n$ to 1.
- Observation has high leverage if h_{ii} is 2 or 3 times $h = (k + 1)/n$, where $k + 1$ is number of coefficients (inc. the constant term).
- As $h_{ii} \rightarrow 1$, x_i may dominate model.

Influence measures — DFBETAS

- An observation that combines “outlyingness” with high leverage exerts an **influence** on the estimated regression coefficients
- If such an observation is deleted from the analysis, the estimated coefficients change substantially.

dfbeta

$$\text{dfbeta}_{ij} = \beta_{j(-i)} - \beta_j$$

dfbetas

$$\text{dfbetas}_{ij} = \frac{\beta_{j(-i)} - \beta_j}{s_{r(i)} \sqrt{(\mathbf{X}^T \mathbf{X})_{jj}}}$$

β_j slope of regression; $\beta_{j(-i)}$ slope when x_i deleted; $s_{r(i)}$ residual SD when x_i deleted; $(\mathbf{X}^T \mathbf{X})_{jj}$ the RSS

- dfbeta_{ij} assesses the impact on the j th coefficient of deleting the i th observation.
- The dfbeta_{ij} are expressed in the metric of the coefficient.
- A standardised version, dfbetas_{ij} divides dfbeta_{ij} by the standard error of β_j .
- Influential observations have $\text{dfbetas}_{ij} \geq 2/\sqrt{n}$

Influence measures — Cook's distance





Cook's Distance

$$D_i = \frac{e_i^2}{s^2(k+1)} \times \frac{h_i}{1-h_i}$$

where e^2 is the squared residual for x_i ; s^2 is the variance of the residuals; h_i is the hat value for x_i

- One problem with dfbetas_{ij} is that there are so many numbers!
- One for each observation for every β_j (inc. the constant); $n \times (k+1)$.
- D_i is a scale invariant measure of distance between β_j and $\beta_{j(-i)}$.
- The first fraction is a measure of “outlyingness”, the second of leverage.
- $D_i \geq 4/(n-k-1)$ suggested as a cut-off for high values of D_i .

Leverage and influence; example



Influence measures in R

- Several functions extract influence measures from fitted models; see `?influence.measures` for details

```
> head(cooks.distance(mod))
      1      2      3      4      5      6
0.0002843771 0.1987126125 0.2084128586 0.1427614594 0.0092847760 0.0084433941
> head(hatvalues(mod))
      1      2      3      4      5      6
0.14468969 0.11994389 0.09996128 0.09991500 0.08264498 0.05635662
> influence.measures(mod)
Influence measures of
      lm(formula = Age ~ Depth, data = agedat) :

      dfb.1_ dfb.Dpth  dffit cov.r   cook.d    hat inf
1  -0.023418  0.02055 -0.0234 1.257 2.84e-04 0.1447  *
2  -0.652525  0.55579 -0.6541 0.981 1.99e-01 0.1199
3   0.675657 -0.55622  0.6813 0.896 2.08e-01 0.1000
4  -0.546052  0.44948 -0.5506 0.985 1.43e-01 0.0999
....
```

Model selection

- Where we have several candidate covariates for inclusion in a model, we face the problem of selecting a minimal, adequate model
- A minimal, adequate model is one that is complex enough to provide sufficient fit to the observed response but no more complex than is necessary
- Several automated techniques available to help
 - ① Best subsets regression — fit all combination of covariates and choose the best model
 - ② Forward selection — start with no covariates, add the covariate that improves fit most, repeat till no covariate results in significant improvement
 - ③ Backwards elimination — as above but start with all covariates and remove the worst variable as long as the model is not made significantly worse
 - ④ Stepwise regression (forward selection and backward elimination)
- Regardless of method used to select a minimal model, you must be aware that these techniques are not a panacea
- p -values from tests on the selected model do not account for the selection procedure; anti-conservative, too many variables selected

Stepwise regression in R



- Base R contains several functions for stepwise selection
 - ▶ `step()`
 - ▶ `add1()`
 - ▶ `drop1()`
- The latter two allow manual selection by single-term addition (`add1()`) or deletions (`drop1()`)
- `step()` is fully automated
- All do selection using AIC not p values
- Package **MASS** contains
 - ▶ `stepAIC()`
 - ▶ `addterm()`
 - ▶ `droptrem()`
- Uses AIC for selection also
- Practical will contain examples of all of these

Best subset regression

- Identifies the best model of each size
- Can use many statistics but AIC and BIC are commonly used

$$\text{AIC} = -2 \times \log(\mathcal{L}(\beta_i) | \text{data}) + kp$$

- k is a penalty on complexity; AIC: $k = 2$; BIC: $k = \log(n)$
- p is number of parameters in model.
- Best subsets is available in package **leaps**



Subset selection and Shrinkage

- Subset selection often used for 2 reasons:
 - ▶ **Interpretation** — Smaller subset of predictors with strongest effects on response y may be easier to interpret and explain
 - ▶ **Prediction accuracy** — LSQ estimates have low bias but large variance. Can sometimes improve prediction accuracy by **shrinking** the coefficients or setting some to zero. In doing so we sacrifice a bit of bias
- Subset selection leads to a small set of interpretable predictors, with possibly lower error (MSE) than the full model
- Subset selection is a discrete process — predictors are either **in** the model, or **out**
- As a result, this subset model often exhibits high variance, which limits the possible improvement in error
- Shrinkage methods are more continuous than subset selection and do not suffer from high variability to the same degree

Stepwise selection & best-subsets

- Stepwise selection is a combination of forward selection and backward elimination steps
- **Forward selection**: start with no terms in model & sequentially add the variable that best improves the model
- **Backward elimination**: start with the full model & sequentially remove the variable that effects the model least
- **Best-subsets**: consider all possible combinations of models (variables) and select the best model for a range of model sizes or select the best model overall
- Several problems with this however:
 - ▶ selection bias in the estimates of the model coefficients $\hat{\beta}_i$
 - ▶ increased variance of the selected model, and
 - ▶ bias in the standard errors of $\hat{\beta}_i$

Selection bias



- Selection bias occurs in the estimates of the model coefficients $\hat{\beta}_i$ in the selection methods
- This bias arises from the effective imposition of a hard threshold on the **size** of the $\hat{\beta}_i$
- $\hat{\beta}_i = 0$ when i th variable is not selected
- Extreme example from Whittingham et al (2006); 5000 data sets ($n = 10$) drawn from the model:

$$y_i = 1 + 0.8x_i + \varepsilon_i$$

- $\beta = 0.8$, $x_i = 1, 2, \dots, 10$,
 $\varepsilon_i \sim N(\mu = 0, \sigma_i = 1)$
- Selection threshold applied of $\hat{\beta} = 0$ where $p > 0.05$

(top) Distribution of $\hat{\beta}$ when OLS applied to each data set.

(bottom) Distribution of $\hat{\beta}$ when significance threshold applied



Ridge regression

- Ridge regression shrinks the coefficients via imposition of a penalty to restrict their size
- Ridge regression coefficients minimises a penalised RSS

$$\beta_{\text{ridge}} = \underset{\beta}{\operatorname{argmax}} \left\{ \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p x_{ij}\beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

or

$$\beta_{\text{ridge}} = \underset{\beta}{\operatorname{argmax}} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p x_{ij}\beta_j)^2$$

subject to



$$\sum_{j=1}^p \beta_j^2 \leq t$$

Ridge regression

- With collinear variables, $\hat{\beta}_{LSQ}$ are poorly determined and have high variance
- One variable can have a large positive coefficient, counteracted by variable with which it is correlated having a large negative coefficient
- Imposing a constraint on size of the coefficients can alleviate this
- Predictors are standardised before running ridge regression
- Intercept β_0 is not subject to the penalty
- Ridge regression shrinks components in the predictors that have low variance (explain low amounts of the variance in \mathbf{X})

Ridge regression

- Need to select a value for the penalty λ , or for the limit on the size of the coefficients t
- Choose these on basis of GCV criterion or CV
- $\lambda = 0$ gives no shrinkage and $\hat{\beta}_{\text{ridge}} = \hat{\beta}_{\text{LSQ}}$
- Ridge regression applied to the Dipper breeding density data:



The Lasso

- The Lasso is a shrinkage method like the ridge regression but with important differences — namely the Lasso can perform variable selection as well as shrink coefficients
- The lasso finds coefficients $\hat{\beta}_{lasso}$ that minimise a penalised RSS

$$\hat{\beta}_{lasso} = \operatorname{argmax}_{\beta} \left\{ \frac{1}{2} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\}$$

or

$$\beta_{lasso} = \operatorname{argmax}_{\beta} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2$$

subject to



$$\sum_{j=1}^p |\beta_j| \leq t$$

The Lasso

- The predictors are standardised prior to analysis and the intercept is not subjected to the penalty term
- Because of the different penalty, if t is sufficiently small (or λ sufficiently large) some of the $\hat{\beta}_{lasso}$ can be shrunk to 0
- This has the effect of selecting those variables with zero coefficients out of the model
- Optimal values for t or λ are chosen using GCV or CV to find those that minimise the prediction error
- Unlike ridge regression, the lasso doesn't penalise sets of low variance or correlated variables to the same extent, however...
- It does do feature selection for us

The Lasso

- Lasso applied to Dipper density data
- Minimum CV error at $\lambda = 0.276$, simpler model within 1 standard error at $\lambda = 0.581$
- 4 predictors have positive coefficients at best model, 3 at the model with 1 standard error
- Gradient (0.019), Stonefly (0.0028), Caddis (0.0004)



The Elastic Net

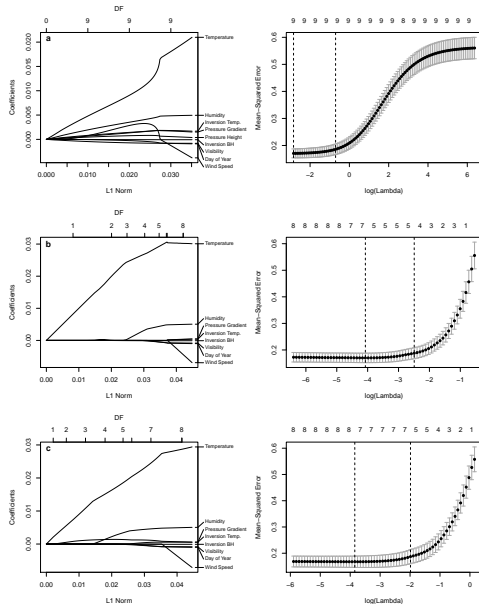
- Ridge regression shrinks all coefficients, proportionally, whilst the Lasso transforms each coefficient by constant factor λ and truncates at zero
- Ridge regression shrinks together the coefficients of correlated data, whilst the Lasso can select or remove coefficients from the model
- Useful if these two properties could be combined
- This is what the **Elastic Net** penalty does
- Find coefficients $\hat{\beta}_{\text{elastic}}$ that minimise the penalised RSS with penalty

$$\lambda \sum_{j=1}^k (\alpha \beta_j^2 + (1 - \alpha) |\beta_j|)$$

- α controls the relative weighting of the ridge-like and lasso-like properties
- Find optimal values of λ and α via a grid search over the parameters using CV and 1se rule

Comparison of shrinkage methods: Ozone data

- Various shrinkage methods applied to predict Ozone concentration using climatic variables
- Left panels show full regularisation paths of $\hat{\beta}_i$ for (a) ridge, (b) lasso, and (c) elastic net
- Right panels show k -fold CV errors for increasing (left to right) penalty
- Dashed vertical lines indicate best model (lowest CV error) and the smallest model within 1 standard error of the best model (right-most dashed line)
- Notice how ridge regression does not perform selection but shrinks correlated variables (Temperature & Wind Speed)
- Lasso performs selection; note difference in paths for Temperature & Wind Speed
- Elastic net ($\alpha = 0.5$) combines both; most similar here to Lasso



Degrees of freedom for shrinkage models

- The degrees of freedom used in finding the fitted values $\text{df}(\hat{y})$ is an important of model complexity
- If we *a priori* present a set of k predictors to linear regression, then that model uses $k + 1$ (for the intercept) $\text{df}(\hat{y})$
- If we do best subset regression, software assumes we have used k df but really we used many more than k
- What about techniques like the lasso and ridge regression?
- Effective degrees of freedom given by

$$\text{df}(\hat{y}) = \frac{1}{\sigma^2} \sum_{i=1}^n \text{Cov}(\hat{y}_i, y_i)$$

- The harder we try to fit the response y_i , the larger their covariance with the fitted values and therefore the more degrees of freedom we have used

Degrees of freedom for shrinkage models

- Effective degrees of freedom given by

$$\text{df}(\hat{y}) = \frac{1}{\sigma^2} \sum_{i=1}^n \text{Cov}(\hat{y}_i, y_i)$$

- The harder we try to fit the response y_i , the larger their covariance with the fitted values and therefore the more degrees of freedom we have used
- This equation works for ordinary regression (it will give k degrees of freedom)
- It works for ridge regression and for the lasso
- In theory this should also work for best subsets regression, but we don't have a closed form equation for estimating $\text{df}(\hat{y})$ in that case
- Highlights the problem of determining the real $\text{df}(\hat{y})$ used if we do best subsets or forward selection / backwards elimination

Multicollinearity redux - VIF

- Variance inflation factor (VIF; V) is related to the sampling variance of a regression coefficient

$$\hat{V}(\hat{\beta}_j) = \frac{s^2}{(n-1)s_j^2} \times \frac{1}{1-R_j^2}$$

where s^2 is estimate error variance, s_j^2 is sample variance of j th covariate

- $\text{VIF} = \frac{1}{1-R_j^2}$ is the **variance-inflation factor** and is a function of the multiple correlation R_j from regression of j th covariate on the other covariates
- $\sqrt{\text{VIF}}$ is a measure of by how much the confidence interval for $\hat{\beta}_j$ is expanded relative to the case where uncorrelated data are used
- $\text{VIF} > \sim 10$ then a covariate is largely explain by other covariates in the model

Multicollinearity redux

- Ridge regression and the lasso estimate biased coefficients
- We accept this extra bias because we attempt to offset the increased variance that complex models and correlated covariates causes
- None of the approaches we talked about is universally a panacea or solution to collinearity
- The real solution is to collect new data so that variables aren't collinear
- Biased estimation methods *may* cause problems worse than collinearity!
- Really, does collinearity actually matter? If we estimate $\hat{\beta}_j$ with sufficient precision then collinearity doesn't matter
- If we can't achieve sufficient precision because of collinearity, this knowledge is only useful if we can redesign the study and collect uncorrelated data
- Think (!) about which terms you introduce to a model

Selected texts

- Fox, J (2008) *Applied regression analysis and generalized linear models*. Sage. (Chapter 13)
- Hastie, T., Tibshirani, R., & Friedman, J. (2010) *The elements of statistical learning*. 2nd Edition. Springer. (Chapter 3). Available from: www.stanford.edu/~hastie/pub.htm
- Whittingham, M.J. et al (2006) Why do we still use stepwise modelling in ecology and behaviour? *Journal of Animal Ecology* **75**:1182–1189
- Murtaugh, P.A. (2009) Performance of several variable-selection methods applied to real ecological data. *Ecology Letters* **12**:1061–1068
- Dahlgren, J.P. (2010) Alternative regression methods are not considered in Murtaugh (2009) or by ecologists in general. *Ecology Letters* **13**:E7–E9
- Simpson & Birks (2012) *Statistical learning in palaeolimnology*. In Birks, H.J.B, Lotter, A.F. Juggins S., and Smol, J.P. (Eds) Tracking Environmental Change Using Lake Sediments, Volume 5: Data Handling and Numerical Techniques. Springer, Dordrecht.

Generalised Linear Models

- Generalised linear models (GLMs) are a synthesis and extension of linear regression plus Poisson, logistic and other regression models
- GLMs extend the types of data and error distributions that can be modelled beyond the Gaussian data of linear regression
- With GLMs we can model count data, binary/presence absence data, and concentration data where the response variable is not continuous.
- Such data have different mean-variance relationships and we would not expect errors to be Gaussian.
- Typical uses of GLMs in ecology are
 - ▶ Poisson GLM for count data
 - ▶ Logistic GLM for presence absence data
 - ▶ Gamma GLM for non-negative or positive continuous data
- GLMs can handle many problems that appear non-linear
- Not necessary to transform data as this is handled as part of the GLM process

Structure of a GLM

A GLM consists of three components, chosen/specified by the user

- 1 A random component, specifying the conditional distribution of the response Y_i given the values of the explanatory data. **Error Function**
- 2 A **Linear Predictor** η — the linear function of regressors

$$\eta_i = \alpha + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_k X_{ik}$$

The X_{ij} are prescribed functions of the explanatory variables and can be transformed variables, dummy variables, polynomial terms, interactions etc.

- 3 A smooth and invertible **Link Function** $g(\cdot)$, which transforms the expectation of the response $\mu_i \equiv E(Y_i)$ to the linear predictor

$$g(\mu_i) = \eta_i = \alpha + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_k X_{ik}$$

As $g(\cdot)$ is invertible, we can write

$$\mu_i = g^{-1}(\eta_i) = g^{-1}(\alpha + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_k X_{ik})$$

GLM Error Function

- Originally GLMs were specified for error distribution functions belonging to the exponential family of probability distributions
- Continuous probability distributions
 - ▶ Normal (linear regression)
 - ▶ Weibull
 - ▶ Gamma (data with constant coefficient of variation)
 - ▶ Exponential (time to death, survival analysis)
 - ▶ Chi-squared
 - ▶ Inverse-Gaussian
- Discrete probability distributions
 - ▶ Poisson (count data)
 - ▶ Binomial (0/1 data, proportions)
 - ▶ Multinomial
 - ▶ Hypergeometric
 - ▶ Pascal
- Choice depends on range of Y_i and on the relationship between the variance and the expectation of Y_i

GLM Error Function

Characteristics of common GLM probability distributions

Probability	Canonical Link	Range of Y_i	Variance function
Gaussian	Identity	$(-\infty, +\infty)$	ϕ
Poisson	Log	$0, 1, 2, \dots, \infty$	μ_i
Binomial	Logit	$\frac{0, 1, \dots, n_i}{n_i}$	$\frac{\mu_i(1-\mu_i)}{n_i}$
Gamma	Inverse	$(0, \infty)$	$\phi\mu_i^2$
Inverse-Gaussian	Inverse-square	$(0, \infty)$	$\phi\mu_i^3$

ϕ is the dispersion parameter; μ_i is the expectation of Y_i . In the binomial family, n_i is the number of trials

Ecologically Error Function

Normal errors rarely adequate in ecology; GLMs offer ecologically meaningful alternatives

- **Poisson** — counts; integers, non-negative, variance increases with mean
- **Binomial** — observed proportions from a total; integers, non-negative, bounded at 0 and 1, variance largest at $\pi = 0.5$
- **Binomial** — presence absence data; discrete values, 0 and 1, models probability of success
- **Gamma** — concentrations; non-negative (strictly positive with log link) real values, variance increases with mean, many zero values and some high values

Logistic regression — Darlingtonia

- Timed censuses at 42 randomly-chosen leaves of the cobra lily (*Darlingtonia californica*)
- Recorded number of wasp visits at 10 of the 42 leaves
- Test hypothesis that the probability of visitation is related to leaf height
- Response is dichotomous variable (0/1)
- A suitable model is the logistic model

$$\pi = \frac{e^{\beta_0 + \beta_1 X_i}}{1 + e^{\beta_0 + \beta_1 X_i}}$$

- The logit transformation produces

$$\log_e \left(\frac{\pi}{1 - \pi} \right) = \beta_0 + \beta_1 X_i$$

- This is the logistic regression and it is a special case of the GLM, with a binomial error distribution and the logit link function

Logistic regression — Darlingtonia

$$\log_e \left(\frac{\pi}{1 - \pi} \right) = \beta_0 + \beta_1 X_i$$

- β_0 is a type of intercept; determines the probability of success ($Y_i = 1$) π where $X = 0$
- If $\beta_0 = 0$ then $\pi = 0.5$
- β_1 is similar to the slope and determines how steeply the fitted logistic curve rises to the maximum value of $\pi = 1$
- Together, β_0 and β_1 specify the range of the X variable over which most of the rise occurs and determine how quickly the probability rises from 0 to 1
- Estimate the model parameters using **Maximum Likelihood**; find parameter values that make the observed data most probable

Logistic regression — Darlingtonia

```
> mod <- glm(visited ~ leafHeight, data = wasp, family = binomial)
> mod

Call:  glm(formula = visited ~ leafHeight, family = binomial, data = wasp)

Coefficients:
(Intercept)  leafHeight
      -7.2930       0.1154

Degrees of Freedom: 41 Total (i.e. Null);  40 Residual
Null Deviance:      46.11
Residual Deviance: 26.96    AIC: 30.96

> ilogit(coef(mod))
(Intercept)  leafHeight
0.0006798556 0.5288181121
```

Logistic regression — Darlingtonia

```
> summary(mod)
```

Call:

```
glm(formula = visited ~ leafHeight, family = binomial, data = wasp)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-2.18274	-0.46820	-0.23897	-0.08519	1.90573

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-7.29295	2.16081	-3.375	0.000738	***
leafHeight	0.11540	0.03655	3.158	0.001591	**


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 46.105 on 41 degrees of freedom
Residual deviance: 26.963 on 40 degrees of freedom
AIC: 30.963

Number of Fisher Scoring iterations: 6

Logistic regression — Darlingtonia



Wald statistics

- z values are Wald statistics, which under the null hypothesis follow a normal distribution
- Tests the null hypothesis that $\beta_i = 0$

$$z = \hat{\beta}_i / \text{SE}(\hat{\beta}_i)$$

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-7.29295	2.16081	-3.375	0.000738	***
leafHeight	0.11540	0.03655	3.158	0.001591	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Deviance

- In least squares we have the residual sum of squares as the measure of lack of fitted
- In GLMs, **deviance** plays the same role
- Deviance is defined as twice the log likelihood of the observed data under the current model
- Deviance is defined relative to an arbitrary constant — only **differences** of deviances have any meaning
- Differences in deviances are also known as ratios of likelihoods
- An alternative to the Wald tests are deviance ratio or likelihood ratio tests

$$F = \frac{(D_a - D_b)/(\text{df}_a - \text{df}_b)}{D_b/\text{df}_b}$$

- D_j deviance of model, where we test if model A is a significant improvement over model B; df_k are the degrees of freedom of the respective model

A Gamma GLM — simple age-depth modelling

- Radiocarbon age estimates from depths within a peat bog (Brew & Maddy, 1995, QRA Technical Guide No. 5)
- Estimate accumulation rate; assumption here is linear accumulation
- Uncertainty or error is greater at depth; mean variance relationship
- Here, fit mid-depth & mid-calibrated age points

	upperDepth	lowerDepth	ageBP	ageError	calUpper	calLower
SRR-4556	20	22.00	355	35	509	307
SRR-4557	26	28.00	465	35	542	480
SRR-4558	32	34.00	635	35	671	545
SRR-4559	38	40.00	740	35	732	666
SRR-4560	44	46.00	865	35	916	691
SRR-4561	50	52.50	870	35	918	692
SRR-4562	56	58.00	985	35	967	795
SRR-4563	100	108.00	1270	35	1284	1097
SRR-4564	200	207.00	2575	35	2761	2558
SRR-4565	260	268.00	3370	35	3697	3487
SRR-4566	400	407.00	4675	35	5563	5306
SRR-4567	493	500.00	5315	35	6263	5955

A Gamma GLM — simple age-depth modelling

```
> plot(calMid ~ midDepth, data = peat,  
+      pch = 21, bg = "black")  
> m2 <- glm(calMid ~ midDepth, data = peat,  
+          family = Gamma(link = "identity"))  
> summary(m2)
```

```
Call:  
glm(formula = calMid ~ midDepth,  
     family = Gamma(link = "identity"), data = peat)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.196221	-0.012606	-0.001604	0.050645	0.092314

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	181.0393	26.0842	6.941	3.99e-05 ***
midDepth	12.2807	0.5025	24.441	3.00e-10 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Gamma family taken to be 0.005924447)

Null deviance: 10.439047 on 11 degrees of freedom
Residual deviance: 0.063394 on 10 degrees of freedom
AIC: 148.83

Number of Fisher Scoring iterations: 4

- Linear relationship
- Error increases with mean & Gamma errors
- Identity link function maintains linearity



A Gamma GLM — simple age-depth modelling

```
> anova(m2, test = "F")  
Analysis of Deviance Table
```

```
Model: Gamma, link: identity
```


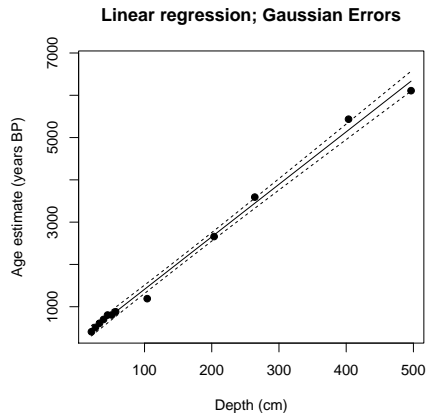
```
Response: calMid
```

```
Terms added sequentially (first to last)
```

	Df	Deviance	Resid. Df	Resid. Dev	F	Pr(>F)
NULL			11	10.4390		
midDepth	1	10.376	10	0.0634	1751.3	1.455e-12 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

A Gamma GLM — simple age-depth modelling




Scatterplots and local relationships

- In scatter plots, it is not always easy to see the form of the relationship between variables
- Ozone concentration tends to decrease as wind speed increases
- But it is difficult to judge whether this relationship is linear or non-linear



Scatterplots and local relationships


- In scatter plots, it is not always easy to see the form of the relationship between variables
- Ozone concentration tends to decrease as wind speed increases
- But it is difficult to judge whether this relationship is linear or non-linear
- Smoothers model the local patterns in a bivariate scatter plot to illustrate the trends or patterns in the data
- They determine the pattern from the data themselves rather than from an *a priori* defined model
- Loess (or Lowess) is one such smoothing technique



Lowess — Locally weighted regression

Locally weighted regression scatterplot smoother





- Decide how smooth relationship should be (span or size of bandwidth window)
- For target point assign weights to observations based on adjacency to target point
- Fit linear (polynomial) regression to predict target using weighted least squares; repeat
- Compute residuals & estimate robustness weights based on residuals; well-fitted points have high weight
- Repeat Loess procedure with new weights based on robustness and distance weights



Try different span and degree of polynomial to optimise fit





Lowess — Locally weighted regression

- Two key choices in Loess
- α is the span or bandwidth parameter, controls the size of the window about the target observation
- Observation outside the window have 0 weight
- Larger the window the more global the fit — smooth
- The smaller the window the more local the fit — rough
- λ is the degree of polynomial using the the weighted least squares
- $\lambda = 1$ is a linear fit, $\lambda = 2$ is a quadratic fit



Lowess — Locally weighted regression

- Two key choices in Loess
- α is the span or bandwidth parameter, controls the size of the window about the target observation
- Observation outside the window have 0 weight
- Larger the window the more global the fit — smooth
- The smaller the window the more local the fit — rough
- λ is the degree of polynomial using the the weighted least squares
- $\lambda = 1$ is a linear fit, $\lambda = 2$ is a quadratic fit



Lowess — Locally weighted regression

“In any specific application of LOESS, the choice of the two parameters α and λ must be based upon a combination of judgement and trial and error. There is no substitute for the latter”

Cleveland (1993) Visualising Data. AT&T Bell Laboratories

- CV can be used to optimise α and λ to guard against overfitting the local pattern by producing too rough a smoother or missing local pattern by producing too smooth a smoother
- However, there are techniques with better properties such as splines that have fewer parameters to choose and which are more widely used
- Loess is perhaps most useful as an exploratory technique as part of EDA
- Cleveland, W.S. (1979) J. Amer. Stat. Assoc. **74**, 829–836
- Cleveland, W.S. (1994) The Elements of Graphing Data. AT&T Bell Laboratories
- Efron, B & Tibshirani, R (1981) Science **253**, 390–395


Splines

- Splines are mathematical functions that take their name from the flexible strips of materials draughtsmen used to draw curves
- A simple spline would just connect the dots, joining each observation to the next — minimal error but rough
- Impose a penalty (λ) on the degree of roughness, so fitting the spline balances the error (lack of fit to the data) with the complexity (roughness) of the spline — **smoothing spline**
- Smoothing splines useful alternative to Lowess for EDA and scatterplot smoothing
- Smoothing splines consist of a series of cubic polynomials over intervals of the data, with intervals defined by knots — piecewise cubic polynomial which is continuous as are it's first a second derivatives

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3$$

Splines

- All the smooths covered here are based on *splines*. Here's the basic idea ...




- Mathematically the red curve is the *function* minimizing

$$\sum_i (y_i - f(x_i))^2 + \lambda \int f''(x)^2 dx.$$

Source: Simon Wood

Splines have variable stiffness

- Varying the flexibility of the strip (i.e. varying λ) changes the *spline function* curve.





- But irrespective of λ the spline functions always have the same basis.

Splines

- **Regression splines** are an alternative type of spline more commonly found in statistical techniques (GAMs)
- In smoothing splines, the observations are the knots and the smoothness is controlled by roughness penalty λ
- In regression splines, a smaller set of knots is chosen across range of the data and cubic polynomials are fitted to the intervals defined by the knots
- As a result, in regression splines the number of knots controls the smoothness of the fitted function
- Once the knots are chosen, regression splines are arguably a parametric approach as we only need to determine the coefficients for the parametric cubic polynomials fitted to each interval
- Regression splines more closely link with formal statistical modelling — can include spline terms in linear regression models and use least squares to estimate parameters



Basis functions — cubic regression splines

- Cubic regression spline basis function takes value 1 at one knot and 0 at others
- j th basis function is multiplied by it's coefficient β_j and then each of these curves is summed at the values of x to yield the smooth curve



Basis functions — cyclic cubic regression splines

- Where x represents a cyclic variable, want ends points of spline to join up smoothly
- Additional constraints on basis functions; second derivatives must match at $f(x_1)$ and $f(x_k)$ (i.e. knots at end points)



Generalised Additive Models

- Generalised Additive Models (GAMs) are a semi-parametric extension of the GLM

$$g(\mu_i) = \eta_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \cdots + \beta_k x_{ki}$$

- GLM requires an *a priori* statistical model
- What if the response can not be well modelled using the available model forms?
- Despite their flexibility, GLMs may not be flexible enough to approximate the true response adequately
- GLMs are model driven
- GAMs include smooth terms of one or more predictors rather than parametric terms
- The form of the smoothers is derived from the data — GAMs are **data driven**

Generalised Additive Models

- Generalised Additive Models (GAMs) for a single covariate has the form

$$g(\mu_i) = \eta_i = \beta_0 + f_1(x_{1i})$$

- The models are additive as all we assume is that the model terms combine in an additive manner to produce the fitted values of the response
- A GAM consisting of smooth terms for several variables has the form

$$g(\mu_i) = \eta_i = \beta_0 + f_1(x_{1i}) + f_2(x_{2i}) + \cdots f_k(x_{ki}) = \beta_0 + \sum_{k=1}^m f_k(x_{ki})$$

- The smooth functions can one of many types of smoother — splines
- Need to specify the type of smoother and complexity of each smoother
- The degree of smoothing for each smooth term can be estimated as part of the model fitting

GAM — Strontium isotope ratios

```
> require(mgcv)
> m <- gam(strontium.ratio ~ s(age), data = fossil,
+         method = "REML")
> summary(m)
```

Family: gaussian
Link function: identity

Formula:
strontium.ratio ~ s(age)

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.074e-01	2.551e-06	277241	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(age)	8.244	8.84	88	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.881 Deviance explained = 89%
REML score = -930.01 Scale est. = 6.9006e-10 n = 106

GAM fit; TPRS; REML smoothness selection



Generalised Additive Models

- In all other respects, GAMs are just like GLMs (link functions, error distributions, etc)
- Using modern methods, the degree of smoothing can be determined alongside the other model parameters using ML
- Interactions can be modelled using a smooth function of two or more variables

$$g(\mu_i) = \eta_i = \beta_0 + f_1(x_{1i}, x_{2i})$$

- In above, thin plate splines impose same degree of smoothing on both variables, tensor product smooths allow for different amounts of smoothing
- Cyclic variables may be modelled using a cyclic smoother; the end points of the smoother are forced to match with no discontinuity

Selected texts

- Wood, S.N. (2006) Generalised additive models; and introduction with R. Chapman & Hall/CRC
- Ruppert, Wand, & Carroll (2003) Semiparametric regression. Cambridge University Press
- Faraway (2006) Extending the linear model with R; generalized linear, mixed effects and nonparametric regression models. Chapman & Hall/CRC
- Zuur, Ieno, Walker, Saveliev, & Smith (2009) Mixed effects models and extensions in ecology with R. Springer

Miscellaneous R commands for working with models

- It is recommended to use extractor functions for the model object
- Common extractor and utility functions are:
 - ▶ `coef()`: model coefficients
 - ▶ `fitted()`: fitted values
 - ▶ `resid()`: model residuals
 - ▶ `vcov()`: variance-covariance matrix of main model parameters
 - ▶ `predict()`: predict from model
 - ▶ `extractAIC()`, `AIC()`: AIC of model
 - ▶ `logLik()`: log likelihood of fitted model
 - ▶ `print()`: quick textual display of object
 - ▶ `summary()`: longer textual display of object
 - ▶ `plot()`: plot the model diagnostics
 - ▶ `add1()`, `drop1()`: add/delete single terms
 - ▶ `update()`: refit the model with changes to formula
 - ▶ `anova()`: partition variance amongst terms or compare models